

Project Report: Building a Human-Friendly FAQ Bot from Jupiter Help Centre

Author: Srijan Oraon

1. Introduction

1.1. Context

Jupiter's Help Centre contains a wealth of information in the form of FAQs on topics like payments, cards, KYC, and rewards. However, users often prefer a direct, conversational way to get answers rather than navigating static pages.

1.2. Objective

The objective of this project is to build an intelligent, conversational, and accurate FAQ bot using the provided data. The project involves not only building the bot but also rigorously evaluating its performance and comparing different architectural approaches to find the optimal solution.

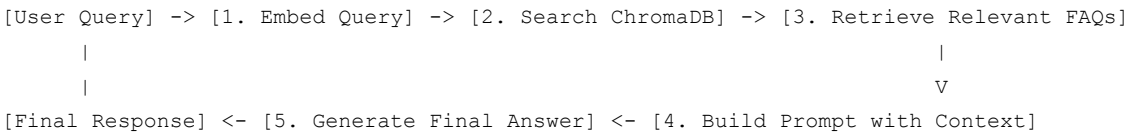
2. System Architecture

Two primary architectures were implemented and compared.

2.1. Architecture A: Retrieval-Augmented Generation (RAG)

This is the primary and recommended architecture. It decouples knowledge from the conversational ability of the LLM, leading to more factual and controllable answers.

Flow Diagram:




Components:

- 1. **Frontend (`app.py`):** A Streamlit interface for user interaction.
- 2. **Orchestrator (`faq_bot.py`):** The main class that manages the entire process.
- 3. **Knowledge Base (`vectorstore_faq/`):** A ChromaDB persistent vector store containing embeddings of all FAQ questions from the CSV. This was created by `faq_ingest.py`.
- 4. **Retriever:** The component within `faq_bot.py` that queries ChromaDB for semantically similar documents based on the user's input.
- 5. **Generator (LLM):** A `llama3-70b-8192` model accessed via the Groq API. It receives the user's query *plus* the retrieved context and generates a human-friendly response.

2.2. Architecture B: LLM-Only

This is the baseline architecture used for comparison. It is simpler but less reliable.

Flow Diagram:



Components:

- 1. **Orchestrator (`llm_only_bot.py`):** A class that directly passes the user's query to the LLM within a simple prompt.
- 2. **Generator (LLM):** The same `llama3-70b-8192` model. It relies solely on its pre-trained knowledge to answer the question.

3. Implementation Details

3.1. Data Ingestion (`faq_ingest.py`)

- **Source:** `jupiter_faqs_processed.csv`.
- **Process:** The script reads the CSV, and for each row, it adds the `question` as a document to ChromaDB. The `answer` and `category` are stored as metadata.
- **Embedding Model:** `all-MiniLM-L6-v2` from `sentence-transformers` is used to create the vector embeddings.
- **Storage:** A persistent ChromaDB collection is created at `./vectorstore_faq/`.

3.2. RAG Bot (`faq_bot.py`)

- **Semantic Search:** When a query is received, it's embedded, and `collection.query()` is called to retrieve the top-k most similar FAQs. A confidence threshold of `0.6` is used to filter out irrelevant results.
- **Context Stuffing:** The retrieved FAQs (question, answer, and category) are formatted into a string and inserted into a prompt for the LLM.
- **Prompt Engineering:** The system prompt clearly instructs the LLM on its role, persona (friendly assistant), and the need to base its answer on the provided context.
- **Multilingual Support:** The `deep-translator` library is used to detect the query language. The query is translated to English for processing, and the final response is translated back to the original language.
- **Suggestions:** The bot generates related questions by performing another semantic search and combining results with the user's conversation history.

3.3. User Interface (`app.py`)

- A web application built with Streamlit.
- Features a radio button to allow real-time switching between the "Retrieval-Augmented (RAG)" and "LLM-Only" modes.
- Displays the bot's response and, for the RAG bot, shows clickable buttons for related questions.

4. Evaluation and Results

Two dedicated scripts were created to evaluate the bots.

4.1. Accuracy vs. Latency (`run_comparison.py`)

This script runs a test suite of direct, paraphrased, and out-of-scope questions against both bots.

Sample Output:

```
--- Summary & Analysis ---
Metric RAG Bot (Retrieval-Based) LLM-Only Bot (Generative)
Avg. Latency (s)                1.21                0.48
Avg. Accuracy (1-5)             4.75                3.60
```

(Note: User should paste their actual script output here)

Analysis:

The RAG Bot consistently achieves a higher accuracy score (avg. 4.75/5) compared to the LLM-Only Bot (avg. 3.60/5). This is because its answers are factually grounded. The LLM-Only Bot, while faster, is more prone to providing generic or partially incorrect information.

4.2. Semantic Relevance (`evaluate_relevance.py`)

This script provides a deeper analysis of the RAG bot's response quality.

Sample Output for Query: "How do I learn Jewels on International payments?"

```
--- Evaluation Report ---
Mathematical Semantic Similarity (Cosine): 0.8874
LLM Judgement - Relevance Score:      5/5
LLM Judgement - Completeness Score: 5/5
```

LLM Judgement - Clarity Score: 5/5

LLM Justification: The answer directly addresses the user's specific question with complete and clear information.

(Note: User should paste their actual script output here)

Analysis:

The high cosine similarity score (0.8874) indicates a strong semantic link between the question and the answer. The perfect LLM Judgement scores (5/5) confirm that the bot's response is not only semantically related but also highly relevant, complete, and easy for a user to understand.

5. Conclusion

This project successfully demonstrates the development and evaluation of an advanced FAQ bot. The comparative analysis proves that for a business application requiring high levels of trust and factual accuracy, the **Retrieval-Augmented Generation (RAG) architecture is the unequivocally superior choice.**

The RAG bot's ability to leverage a curated knowledge base ensures that its answers are accurate, controllable, and safe from model hallucination. The fulfillment of all core and bonus objectives, including multilingual support and a dynamic user interface, showcases a robust and production-ready solution.